

UK Heat Split: data sources and estimation methodology

Version 8 · Causeway Energies · August 2026

This revision supersedes the v7.1 methodology. Sections 4, 5, 6.1, 10.3, 10.5 and 12 are substantially rewritten; Sections 4.6, 5.4, 6.4 and 13 are new. A change log against v7.1 is given in Section 14.

1. Conventions

1.1 Sourced quantities and author estimates

Every figure on the site is either traceable to a published source or is a Causeway Energies estimate. Estimates are marked with a dagger (†) wherever they appear, both in this document and on the site itself, and the contact address for challenge is published beside them. The distinction is maintained even where an estimate is well founded: the dagger denotes provenance, not doubt.

1.2 Purchased energy and delivered service

Two quantities are reported throughout and must not be conflated.

Purchased energy is what was bought: gas at the meter, electricity at the grid supply point, oil delivered. It is what appears on a bill.

Delivered service is what the building received: useful heat after combustion losses, useful cooling after the refrigeration cycle. For heat pumps delivered service exceeds purchased energy, because the difference is lifted from the environment rather than bought.

The energy-balance panel shows both, and the gap between them is the sum of combustion losses (a loss) and ambient heat harvested (a gain). These are opposite in sign and must not be netted into a single “efficiency”.

4. Cooling: the equipped fleet (observed)

Substantially revised at v8. The previous version described a latent-slope extrapolation and a saturation test that the data did not support; both have been withdrawn. See Section 14.

4.1 Data source and reconstruction

(unchanged from v7.1)

Cooling delivered by installed plant is inferred from national electricity demand. Half-hourly demand is taken from the National Energy System Operator (NESO) historic demand data and rolling update files [16]. Metered demand understates true underlying demand by the quantity supplied behind the meter, principally embedded solar photovoltaics, which is largest precisely

on the hot, bright days when cooling load peaks. Underlying demand is therefore reconstructed by adding NESO estimates of embedded solar and wind generation to metered demand before analysis. The correction is material: embedded solar alone reaches of order 10–16 GW at midday in high summer against a metered demand of order 20–30 GW. Future-dated rows present in the update file, and days with other than 48 half-hourly values, are removed.

4.2 Within-class centring

(unchanged from v7.1, but see 4.5 for what it costs)

Summer electricity demand carries a strong non-cooling structure: it falls during the August holiday period and differs between weekdays and weekends. Binning raw demand against CDD conflates this structure with the cooling response and can yield a spurious negative uplift. To isolate the cooling signal, each daily underlying-demand value d_i is expressed as an anomaly about the mean of its (calendar-month $\times$ weekend/weekday) class:

$$a_i = d_i - \bar{d}(m_i, w_i)$$

where $\bar{d}(m, w)$ is the mean underlying demand over all days sharing month m and weekend/weekday status w . The operation is the within transformation of standard panel econometrics — a two-way fixed-effects demeaning on calendar month and weekend/weekday class [31] — applied here to isolate a weather signal rather than to estimate a panel coefficient. It is validated against synthetic data with a known injected slope (Section 4.4).

The term “demeaning” is avoided in public-facing copy in favour of “within-class centring”, to prevent confusion with the ordinary meaning of the word.

4.3 Response curve: a continuous convex fit

Changed at v8. The response was previously binned by CDD and read as a step function, with a linear slope fitted to the low-CDD bins and extrapolated to define “latent” demand. Two faults were found.

The binning imposed a ceiling. The top bin held every day from CDD 4 to CDD 7.9 at a single value — fifteen days a year served the same figure whatever the weather.

The weekly headline kept the binned estimator until 17 August 2026. The continuous fit was introduced on 11 August but computed thirty lines *below* the weekly accumulation, which still read the bin lookup — so the same metric had two implementations: the daily chart was continuous and the weekly Tier 1 bar was the old shelf. On a representative curve the binned weekly figure **understates a hot week by about 80%** (256 GWh against 468) and overstates a mild one, because bin 1 spans CDD 0 to 1 at one value. The fit is now computed once, above both, and there is a single estimator. Found in peer review. The response is now fitted continuously as

$$E(c) = k \cdot c^2$$

where c is cooling degree days on the daily mean of population-weighted outdoor air above the base in Section 2.3, and k is fitted by least squares through the observed bin means. The quadratic form is not assumed: it is what an undemeaned regression independently selects

(Section 4.6), with an exponent of 2.0 to two significant figures and R^2 0.96 against the bin means.

The convexity is the extensive margin — each building is a step function and the fleet is the integral of those steps — on three independent grounds. The response *steepens* rather than saturating, which is the opposite of a capacity constraint. Dividing out the falling COP leaves the bend almost intact, so it is load rather than plant. And adding moist-air enthalpy days on the full record moves the quadratic coefficient by **3%**, where a latent-load explanation requires moisture to absorb a material part of it (Section 4.8). That third test was capable of falsifying the reading and did not. Removing the ceiling raises the annual total by about 26%. k is published in the payload as `k_quadratic` so the fit can be reproduced.

The response is convex, not saturating. The previous version tested for capacity saturation by extrapolating the low-CDD slope and looking for the observed curve to fall below it. The measured response does the opposite: it steps about +3 GWh per CDD across the cool bins against +17 across the warm ones. A linear extrapolation from the shallow end therefore sits *below* the observed curve at high CDD, and the difference it was meant to measure is structurally zero. The saturation tier has been withdrawn (Section 5.3), and the latent-slope line has been removed from the published chart.

4.4 Validation

(unchanged from v7.1)

The centring estimator was validated before deployment on synthetic data in which a known cooling slope was combined with an injected holiday depression of comparable magnitude to that seen in real August demand (Figure 3). Naive binning of the synthetic series reproduced the spurious negative-uplift artefact; within-class centring recovered the injected slope to within sampling error.

4.5 Why the observed and anchored figures differ

Rewritten at v8. The previous version reported an “order-of-magnitude discrepancy” between the observed response and the ECUK-anchored figure and left it unreconciled. The cause has since been identified and is a property of the estimator, not an error in either quantity.

Within-class centring removes the level along with the confound. Subtracting the (month $\times$ weekend-class) mean removes the mean itself. Cooling that runs steadily through July is contained in the July mean and is therefore invisible to the centred estimator; only day-to-day variation about that mean survives. The observed response is consequently a **lower bound by construction**, not an independent estimate of the total.

Quantitatively, the published Tier 1 recovers about **33%** of the ECUK cooling anchor - 4,681 GWh thermal on the trailing year against 14,352. **Corrected at v8.47 from a stated “about a seventh”.** That figure was the raw within-class centred curve before the binned ceiling was replaced by the continuous quadratic on 11 August, which raised the total 26%, and before later corrections. The published series had moved and the copy had not. Modelling the baseline explicitly instead — a constant, a day-type term, holiday terms, a heating limb and a convex cooling limb, fitted on **undemeaned** daily data — retains the intercept and recovers **74%**

(Figure 5). Two estimators, one anchor, and the difference between them explained by the operation that separates them.




Figure 5

Figure 5. Two estimators against the ECUK cooling anchor. Within-class centring removes the level along with the confound and recovers about a seventh of it; modelling the baseline explicitly on undemeaned daily data retains the intercept and recovers 74 per cent.

This also resolves the tier arithmetic: the measured response and the anchored figure are not additive. The measured tier is a **subset** of the anchor, being the weather-varying part of the same equipped stock's consumption. The tiers that lie *outside* the anchor are those the anchor cannot see — buildings that want cooling and do not have it (Section 5), because a consumption survey cannot count a building that consumes nothing.

4.6 The undemeaned baseline model

New at v8. Because the centred estimator cannot recover a level, a second estimator is run offline for reconciliation. Daily GB demand, with embedded generation added back and no centring applied, is regressed on

$$D = \beta_0 + \beta_1 \cdot \text{weekend} + \beta_2 \cdot \text{HDD}(15.5) + \beta_3 \cdot \max(0, \bar{T} - T_c)^2 + \sum \beta_j \cdot \text{holiday}_j$$

where $\bar{T}$ is the population-weighted daily mean and T_c the cooling balance point, scanned over 11–20 °C. Because the baseline is modelled rather than removed, the intercept survives and T_c is identified.

Superseded 17 August 2026. The originally published result — $T_c = 14.5$ °C, 3.54 TWh, 74% of the anchor — rested on the double-counted solar of Section 4.7 and is withdrawn. Two corrected runs on partial records gave 24–37%, which was too wide to publish.

On the corrected series with embedded solar carried as a control and the full 800-day record, the estimator reconciles far better than it ever has:

Specification	Balance point	Cooling	Share of anchor	R ²
Base: constant, weekend, HDD, convex cooling limb, solar	12.0 °C	3.59 TWh	75%	0.856
Full: plus holidays, Christmas, August dip, trend, moisture	11.0 °C	4.56 TWh	95%	0.918

Carrying solar was the change that mattered. It was previously excluded for entering “with the wrong sign for a lighting effect” — reasoning that rested on the contaminated variable, since the dependent series already contained the embedded solar being tested. Corrected, it enters at -0.91 with the largest incremental R² of any single control, and a negative coefficient is what embedded generation displacing metered demand should produce.

The full specification carries nine regressors and 95% should not be read as a validation. Each term is a documented control with a physical rationale rather than a fitted nuisance parameter, and an earlier mis-specified version of the same model returned 164% of the anchor — which is what overfitting actually looks like. But the honest reading is that the undemeaned estimator now recovers most of the anchor and is no longer obviously biased, not that the two have been reconciled to within 5%.

This estimator is still not used in the published bill or carbon figures, which remain ECUK-anchored. The case for switching is materially stronger than it was, and is deferred to a second full cooling season.

It is now drawn on the site, over the Tier 1 bars (v8.44, relabelled at v8.46). The two are **not the same quantity** and the label was corrected to say so: the bars are the weather-varying slice the centred method can see, and the line is an estimate of the equipped fleet’s *whole* cooling load. The gap between them is therefore informative in itself - it is approximately what within-class centring removes, which is the subject of Section 4.5.

The line is computed live from the published daily temperature series, not hardcoded — but its coefficient is **pinned** to the fit that produced it (UNDEMEANED_FIT in `build.py`: $k = 0.6109$ GWh per °C² above 11.0, fitted 2026-08-17 on 800 days with nine regressors, R² 0.918). A coefficient without its vintage is not reproducible, so the fit date, the record and the resulting share of the anchor are stored alongside it and must be updated together when `analysis/fit_cooling.py` is rerun. This is a first instance of the version pinning discussed in Section 12.

The line is absent, and its legend entry suppressed, on any payload built before v8.44 — an old `data.json` degrades to the two bars rather than showing a broken series.

Adding controls moves the cooling term by less than $\pm 10\%$: holiday terms contribute the most ($\Delta R^2 +0.058$ for the Christmas shutdown, $+0.046$ for bank holidays), a linear trend $+0.018$, and embedded solar $+0.008$. The solar coefficient enters with the wrong sign for a lighting effect and is collinear with temperature; it is not carried.

This estimator is **not** used in the published bill or carbon figures, which remain ECUK-anchored. It is a reconciliation diagnostic, and the case for switching the site onto it is deferred pending a full second cooling season.

4.7 Diurnal separation and the night response

New at v8. Two hourly diagnostics are computed and published, both from the retrospective store.

The diurnal slope profile — WITHDRAWN 17 August 2026. This section previously reported an overnight floor of $0.261 \text{ GW}/^\circ\text{C}$ rising to 0.676 at noon, and offered the $2.6\times$ separation as evidence that a continuously running refrigeration load is distinguishable from daytime comfort cooling.

It was an artefact of a double-counted variable. demand_GW in the retrospective store is already *underlying* demand — NESO national demand plus embedded wind **and embedded solar**. The diagnostic added embedded solar a second time, which manufactures a correlation with temperature in precisely the daylight hours where a cooling signal would appear.

Corrected, the profile is flat: **$0.262 \text{ GW}/^\circ\text{C}$ at the overnight floor and 0.243 at noon**, a ratio of 0.93 against the published 2.59 . There is no daytime separation in this data. The chart and the claim are withdrawn from the site, and Tier 0 is now supported by the bottom-up count alone.

The same error was present in the offline reconciliation script, and was found in peer review rather than by the build gates — a slope diagnostic has no plausible range to gate against.




Figure 6

Figure 6. Summer demand regressed on temperature, hour by hour, centred within year, month, day type and hour. The overnight floor is cooling that runs regardless of the clock; the daytime excess above it is comfort.

The night elbow. Restricting to 01:00–05:00, May to September, and binning by temperature, the response is **not linear**: it is flat below about 15 °C and steep above (Figure 7). The break replicates independently in each summer on record — -0.074 then $+0.300$ GW/°C in 2025, -0.109 then $+0.426$ in 2026 — on 940 night hours.




Figure 7

Figure 7. Summer-night demand against outdoor temperature, binned, by year. The response is flat below about 15 degC and steep above, and the break replicates independently in each summer on record. Point size shows hours per bin.

No level is inferred from either. An earlier attempt to invert the measured slope into an absolute refrigeration load via the COP identity was withdrawn: the flat segment below the break means the line being extrapolated does not exist, and the height of that floor is precisely the quantity aggregate demand cannot see. This is stated because the inverse method is superficially attractive and a reader may otherwise attempt it.

4.8 Tested and rejected

New at v8. Three candidate refinements were tested against the data and are recorded here as negative results.

Humidity — tested, rejected, and the null now properly obtained. Latent load rises with absolute moisture, so moist-air enthalpy was a plausible explanation for the convexity in Section 4.3. Hourly dewpoint was fetched from the same ERA5 archive on identical population weights, converted to moist-air enthalpy against a 22 °C/50% RH reference (43.03 kJ/kg), and entered as an additional regressor.

The result is a null, and it took three attempts to obtain one that could be relied on.

The history is recorded because it shows what the intermediate results were worth:

Run	Basis	Quadratic moves	Verdict
Original	Double-counted solar, no solar control, 400 days	+7%	Null reported, but obtained under a bias running <i>toward</i> null
First rerun	Corrected series, no solar in base, 400 days	-25%	Reversed — but moisture was proxying for the omitted solar term
Second rerun	Solar in base, 400 days	+27%	Reversed back — but moisture terms fitted on half the record
Final	Solar in base, 800 of 800 days	+3%	Null, properly obtained

On the full record with solar carried in the base, **dewpoint adds 0.0006 to R^2 and enthalpy days add 0.0000**, against 0.031 for bank holidays and 0.028 for the Christmas shutdown. Both moisture terms return a balance point of 12.0 °C, identical to the base — they add nothing because there is nothing left for them to explain. The two also now **agree with each other**, as near-identical constructs should; their earlier sign disagreement was an artefact of being fitted on half the record while every other term used all of it.

The site's own stated test is that if adding enthalpy days barely moves the quadratic coefficient, the bend is the extensive margin rather than latent load. It moves **+3%**. The extensive-margin reading therefore stands, and stands on a test that was capable of falsifying it — the two intermediate runs show it moving 25 and 27 points under mis-specification.

Two defects had to be removed before the test meant anything, and both are worth naming because each on its own produced a confident wrong answer: double-counted embedded solar in the dependent variable (Section 4.7), and the omission of solar as a control, which left moisture proxying for it because humid days are cloudier. A third — humidity covering only half the record, from a workflow input that silently overrode the fetch default — made the incremental R^2 values incomparable. `fit_cooling.py` now warns below 95% coverage and the workflow fails outright below 700 days.

Slope-to-level inversion. See Section 4.7.

The saturation tier. See Sections 4.3 and 5.3.

5. Cooling: the comfort deficit (modelled)

Renumbered and rebased at v8.

5.1 Threshold and basis

Changed at v8. The previous version quantified overheating exposure as degree-hours above 26 °C, citing the CIBSE operative-temperature criterion. This was a misapplication: **26 °C in CIBSE TM59 Criterion B is an INDOOR operative temperature** for bedrooms over 1% of night hours. It was applied to *outdoor* air, which is a different quantity and a much rarer event, and the tier was consequently understated by roughly five-fold.

The threshold is now **20 °C on the daily mean of population-weighted outdoor air**, derived from the 2017 Energy Follow-Up Survey indoor/outdoor regression, in which indoor bedroom temperature rises approximately 0.48 °C per 1 °C of outdoor temperature. Cooling degree-hours are computed as

$$CDH_{20} = \sum_d 24 \cdot \max(0, \bar{T}_d - 20)$$

This base is not the 15.5 °C cooling shape base of Section 2.3 and answers a different question. 15.5 °C is the balance point at which comfort cooling plant begins to run in buildings that have it; 20 °C is the outdoor daily mean above which buildings that have none become uncomfortable. Both appear on the site and the distinction is stated wherever they do.

A night-time urban heat island uplift is not in the numbers and would add an estimated 10–20% (†).

5.2 Stock and response

(largely unchanged; scenario values retained)

Three prevalence scenarios are carried and never averaged: a low case of 11% of dwellings from English Housing Survey self-reported overheating [9]; a central case of 25% (†); and a high case of over one-half, from the Climate Change Committee assessment of dwellings at overheating risk [10]. Non-domestic uncooled comfort space is taken as 190 Mm² (†) from BEES floor areas.

Thermal response is 0.2 kWh thermal per degree-hour per dwelling (†), equivalent to an effective cooled-zone UA of about 200 W/K, and 3.0 Wh thermal per m² per degree-hour for non-domestic space (†). These remain the dominant uncertainties.

Unit convention. This tier is a **thermal load computed directly from building physics and never divided by an efficiency**. Tier 1 by contrast is measured electricity multiplied *up* by the fleet EER. The two are reported in the same units and rest on entirely different assumptions; the site states this in the chart legend.

5.3 Tier structure

Changed at v8. The site previously published three tiers. The middle tier — “equipped but saturated”, defined as latent demand minus delivered — has been **removed**, and the former tier 3 renumbered into its place.

The removed tier’s premise was that the equipped fleet saturates at high load. The measured response steepens instead (Section 4.3), so $\max(0, \text{latent} - \text{delivered})$ was structurally zero. Saturation would present as a *concave* response; the measured one is convex.

The computation is retained as a diagnostic and logged each run: a non-zero value would indicate the curve had turned concave, which is the evidence that would reinstate the tier.

The published tiers are therefore:

Tier	Content	Relation to the ECUK anchor
1	Weather-varying response of the equipped fleet	Inside it — about a third (Section 4.5)
2	Buildings with no cooling equipment at all	Outside it — a consumption survey cannot count non-consumption
0	Process cooling: data centres, cold stores, retail and food refrigeration	Outside ECUK U5, which is a services-buildings table. Of order 26 TWh/yr of service (†); named in the legend and deliberately not plotted, being several times the others combined

5.4 Delivery efficiency

Changed at v8. The comparison route efficiencies were previously stated as “an effective COP of order 20” for ground cooling and “a seasonal EER of order 3.5” for air-source, giving a six-to-one ratio.

The air-source figure is now **3.0, with a published range of 2.5–4.0**. The range matters: it is the largest single unverified assumption in the cooling chain, and it propagates one-for-one into every cooling service figure on the site. Legacy oversized plant runs nearer 2; modern commissioned VRF nearer 4.5. BEES excludes fan energy from its space-cooling line, so the boundary here is nearer chiller-plus-condenser than whole-system, which argues for the upper half of the range; UK part-load operation argues for the lower half.

What would resolve it: air-conditioning inspection reports under EPBD carry system type, age and rated efficiency for a large sample of UK non-domestic plant, and are the one dataset capable of producing a fleet-weighted seasonal figure rather than a judgement.

4.9 The cooling floor: four routes tried, one still open

Added at v8.46, and written as a record of what did not work as much as what did. Section 4.5 establishes that the observed cooling response is a lower bound by construction: within-class centring removes the level along with the seasonal confound, so cooling that runs steadily through July is invisible. The question this section addresses is whether that level can be recovered by any other means. Four routes were tried offline. **None has yet produced a figure the site publishes**, and the two that failed are recorded because a reader should be able to see the attempts as well as the results.

A note on the fractions in this section. The centred estimator's share of the anchor has been restated more than once as the estimator changed. The figure that matters is the PUBLISHED Tier 1 against the published anchor: **33%** on the trailing year at v8.47. Earlier statements of 15% or "a seventh" predate the continuous-quadratic fit and should be read as historical.

Route 1 — model the baseline instead of removing it. Section 4.6. Recovers 75–95% of the anchor and is drawn on the site as a comparison line. **It does not answer the question:** year-round cooling is indistinguishable from any other constant load and sits in the intercept alongside everything else. What it recovers is the total, not the split.

Route 2 — separate comfort from process by hour of day. WITHDRAWN. The published version was an artefact of double-counted embedded solar (Section 4.7). Corrected, the diurnal profile is flat: 0.26 GW/degC at 04:00 against 0.24 at noon, a ratio of 0.93 where 2.59 was claimed.

Route 3 — fit a balance point as the vertex of a V. WITHDRAWN as an identification method. Fitting heating and cooling limbs meeting at a vertex is a better functional form than flat-then-steep, and it explains the negative low-temperature slopes as a genuine heating limb. But the vertex is not identified: regressed on each hour's own mean temperature it returns a slope of 0.98 with R2 0.97. The V-fit places its kink about 1.5 degC above the centre of whatever data it is given. The apparent agreement between two summers was two years with similar night temperatures, not two measurements of one quantity.

Route 4 — borrow a level from the gas side. If a building doing both heating and cooling never does both at once, the distribution of temperatures at which buildings stop heating is also the distribution at which they may start cooling - and gas has an absolute level that electricity lacks. Fitting the gas limb with balance points distributed as $N(\mu, \sigma)$ rather than shared gives **mu 15.5 degC, sigma 2.00 degC** on 612 days, with the RSS improving 2.3% against the nested hard-kink model. That improvement is significant: at the observed residual of 169 GWh/day a true kink returns $\sigma \geq 2.0$ in 0 of 60 simulations and an RSS gain above 0.73% in none of 40.

So balance points are genuinely spread, and the site's single 16.5 degC base is roughly where 70% of the stock has already stopped heating rather than where heating ends. But the route gives a *shape* and not a *scale*: the cooled stock is a different population from the heated stock, and leaving heating does not imply entering cooling. `analysis/fit_balance_spread.py`.

Route 5 — the weekday/weekend contrast. STILL OPEN, and the only route that has produced a level. Compare weekday with weekend demand at the same temperature, hour and season. Everything invariant to occupancy cancels - baseload, process refrigeration, the weather itself - leaving load that responds to both temperature and occupancy.

On summer afternoons the gap grows **+0.107 GW/degC** (R2 0.90, six bins), from 2.64 GW at 13.5 degC to 4.27 at 27. Integrated above a 14 degC reference this is **467 GWh/yr measured in the 12-17h band, or 1,866 GWh/yr scaled across the day - 10% to 39% of the anchor.** The range is the honest answer: the band figure counts six hours a day and understates, the scaled figure assumes the daytime rate runs overnight and overstates.

The placebo passes on sign. Repeating the contrast at night in the same season switches the occupancy difference off while leaving the weather unchanged: the night slope is **-0.065 GW/degC** against the day's +0.107. The gap grows with temperature only when buildings are occupied, which a generic weekday/weekend artefact would not produce. The night gap is 0.94 GW and temperature-insensitive - a standing occupancy load that is not cooling.

Three limits, all material. Buildings open at weekends - retail, hospitality, leisure - appear on both sides and cancel, so this is a lower bound on commercial cooling and possibly a substantial one. Other occupancy-linked load that tracks temperature, ventilation in particular, is not separated from cooling. And the night placebo rests on three bins over four degrees, because Britain has few hot nights; it is directional rather than conclusive and should be rerun as the record lengthens. `analysis/fit_weekend_contrast.py`, which fails rather than reporting a figure if the placebo does not hold.

Nothing from routes 4 or 5 is published on the site. They are recorded here so a reviewer can see the state of the question, and so the next person does not repeat routes 2 and 3.

5.5 The UK-to-GB boundary

Rewritten at v8.40 after peer review. ECUK reports UK totals; the site reports Great Britain, because the live gas signal comes from the GB distribution zones and Northern Ireland is not on that system. Every annual therefore passes through a UK-to-GB scaling step.

Until 17 August 2026 that step applied one factor — 0.985 — to every fuel and end use. That figure is a statement about the *gas network*, and it does not transfer to fuels whose NI penetration is nothing like their GB penetration.

It was materially wrong for oil, where NI holds a far larger share of the UK total than its population implies. Two independent routes agree once they are put on the same basis:

- **NIHE and BRE, “Energy consumption in Northern Ireland’s housing stock: 2016”** — BREDEM modelling on the House Condition Survey — give NI domestic primary space heating of 10,700 GWh with oil at 76%, and water heating of 3,500 GWh with oil at 2,700. That is **10.8 TWh of NI oil heat on a 2016 basis**, 21% of the UK’s 51.0 TWh. The same source puts NI domestic heat in total at **14.2 TWh**, which is within 2% of the 14.0 TWh estimate this site was already carrying for the NI panel.
- **The Irish Heat Split sibling** estimates **8.3 TWh** on a trailing twelve-month basis, 16%.

These are not in conflict — they are nine years apart, in the one jurisdiction where oil is retreating fastest. NISRA puts the oil share of NI homes at 61% in 2024/25 against 68% on the 2016 basis; that decline, nine years of fabric and boiler improvement, and the weather difference between a trailing window and the 2016 year close the gap: $10.8 \times 0.90 \times 0.93 \times 0.97 = 8.8$ against 8.3, a 5% residual.

Because the site reports a trailing window, the current-basis figure applies: **NI at about 17% of UK oil heat, so a GB share of 0.83.** Applying the gas-network 0.985 overstated GB oil heat by about 19%.

A first attempt used 0.69 and was revised the same day. It was derived from household counts alone — 61% of 0.80m NI homes against roughly 1.1m GB oil-heated homes — which ignores consumption per household and overstated the correction, NI dwellings being smaller and newer than the GB off-grid stock, which skews rural and large. It is recorded here because a correction that itself needed correcting is worth showing.

Two known biases, opposed: the NI oil share is still falling, which argues for a higher GB share; and NIHE covers **domestic only**, so NI non-domestic oil sits on top of the 10.8, which argues for a lower one.

The shares are now per fuel:

Fuel	GB share	Basis
Gas	0.988	NISRA 36% of NI homes on gas; GB gas-connected ~23m (†)
Oil	0.830	NIHE/BRE 2016 reconciled to a trailing basis; NI ~17% of UK oil heat (†)
Solid	0.893	Low as an NI primary method, higher as secondary (†)
All others	0.972	No NI-specific driver, so the household share (†)

The effect on total GB annual heat is about **−4%**, almost all of it oil. Every figure that divides by total heat moves correspondingly, including the capacity headroom shares in Section 10.5.

All four factors are estimates. The NI household counts and the GB oil-heated stock are both approximate, and no official UK statistic publishes the GB/NI split of oil heat directly. What has changed is that the shares are now *fuel-specific and reasoned* rather than a gas-network artefact applied uniformly; the residual uncertainty on the oil figure remains material and is daggered.

6.1 In-situ efficiencies, performance factors and the ECUK cooling split

Substantially revised at v8.

Combustion

Combustion appliances are derated by in-situ, not nameplate, efficiencies. From 13 August 2026 **space heating and hot water are converted at different efficiencies:**

Fuel	Space heating	Hot water	Basis
Gas	0.835	0.73	In-situ winter [5]; hot water derived below

Fuel	Space heating	Hot water	Basis
Oil	0.82	0.71	As above
Bioenergy	0.70	0.70	Single figure — not covered by SAP Table 4b
Solid	0.55	0.55	As above

Why they differ. A boiler serving only a cylinder fires in short bursts, cools between them, and returns water above the condensing dewpoint — so it loses its latent recovery outright. SAP has carried separate winter and summer seasonal efficiencies for this reason since SAP 2009.

Derivation. The hot-water figures are the mean SAP 2012 Table 4b summer/winter ratio applied to the existing in-situ winter anchor, so the sourced winter values are untouched:

- gas: ratios 64/74, 74/84, 65/74, 75/84 — mean $0.879 \times 0.835 = \mathbf{0.73}$
- oil: ratios 68/80, 72/84, 73/82 — mean $0.866 \times 0.820 = \mathbf{0.71}$

Summer runs 82–89% of winter across the whole of Table 4b. Source: SAP 2012, Table 4b p207 [32].

The summer figure is applied to hot water in every month, not seasonally. SAP is explicit that $\eta_{\text{water}} = \eta_{\text{summer}}$ for all twelve months: the distinction is between end uses, not between seasons.

Two caveats. Table 4b is SAP's *fallback* for boilers absent from the Product Characteristics Database, so it is conservative and skewed to older plant; a PCDB-weighted fleet figure would be better. And Table 4c deducts five points from both figures where a regular boiler has no interlock, which is not applied here.

Effect. The gas line acquires a seasonal shape it did not previously have. The effective gas efficiency, weighted across both end uses, now runs about 0.81 in January and **0.74 in July**, when the load is almost entirely hot water. It was a constant 0.835 in every month before. Useful annual hot water falls 66.2 — 58.8 TWh, so the space-plus-water total falls 1.8% and the electrification limit, which divides by that total, rises by about the same.

Heat pumps

Seasonal performance factors are 2.80 for air-source, from the Energy Systems Catapult Electrification of Heat demonstration [7]; 3.24 for ground-source, from Energy Systems Catapult in-situ monitoring of more than 1,700 installations [6]; and 5.0 for a networked geothermal ambient loop (†), for which no monitored UK fleet exists.

System boundary, stated explicitly at v8.42 after peer review. All three figures are whole-system, on the same basis as the SEPOMO H4 boundary the two monitored anchors come from: **electricity metered at the supply point over heat delivered at the building boundary.** Inside that boundary for the network route are the building heat pumps and their auxiliaries, district ambient-loop circulation, source-side and primary-loop circulation, production and injection well pumping where a doublet is used, controls, and the round-trip loss on the seasonal store — and because the denominator is heat delivered at the building, distribution

heat losses are inside too. Outside it are groundwater treatment where required (not modelled, and would reduce the figure), customer-side emitters, and anything embodied.

This matters because a monitored H4 SPF compared against a plant-only network COP would not be like-for-like, and the comparison would flatter the network throughout. **The parasitic load is the part most likely to be optimistic.** European ambient-loop and aquifer-storage schemes commonly report 5–15% of delivered heat going to circulation; a plant-level 5.6 becomes a whole-system 5.0 at about 10%, which is the arithmetic this figure assumes. At 15% it would be 4.87 and at 5% it would be 5.33.

These annual factors anchor the model; they are not applied as constants in the hourly or daily engines. See Section 10.3.

The ECUK cooling and ventilation split

New at v8, and material. ECUK End Use Table U5 publishes services-sector “cooling and ventilation” as a single electricity line, held here at 10.4 TWh/yr. The site previously split this 50/50 between a weather-flat ventilation component and a weather-shaped cooling component. That split was undocumented and is now sourced.

The Building Energy Efficiency Survey (BEES, DECC/BEIS, fieldwork 2014–15), which is the survey ECUK’s services end-use split is derived from, reports the components separately: **fans 6,030 GWh, space cooling 5,090 GWh and humidification 130 GWh** of electricity. That is **54% ventilation to 46% cooling**. US CBECS 2018, which separates the two natively, gives 18% of commercial electricity to ventilation against 14% to cooling — a 56:44 ratio from a considerably hotter climate, which corroborates the shape.

Ventilation has no coefficient of performance. Fans move air, not heat; there is no thermodynamic multiplication to apply, and fan work dissipates as heat inside the space, making it a cooling *load* rather than a cooling output. Ventilation is therefore reported as its own service at unity and is **excluded from the cooling figure**. This correction reduced published cooling delivered by approximately one third, from 20.8 to 14.35 TWh/yr on the annual anchor.

The cooling shape base

Changed at v8 from 18.0 °C to **15.5 °C** on the daily mean. 15.5 °C is the balance point at which comfort cooling begins in buildings with high internal gains, is the traditional UK degree-day base, and is what CIBSE TM46 applies across all sectors. It is corroborated by the undemeaned regression of Section 4.6, which places the best-fit balance point at 14.5 °C.

The Met Office cooling-degree-day base of 22 °C daily mean is a climate-comparison convention, not a building balance point, and would understate plant operation substantially. It is not used.

Degree days are fetched at bases 15.5, 18.0, 20.0 and 22.0. **CDD is convex in its base and the series cannot be interpolated between them** — averaging two bases overstates the middle.

6.4 The heat-price series

New at v8. A daily series of the commodity cost of delivered heat is published for four routes on two bases.

Wholesale. Gas is the National Gas System Average Price divided by the 0.835 boiler efficiency. The three electric routes are the daily electricity market index divided by **each route's own COP on that day's temperature** (Section 10.3), not by a fixed seasonal factor.

Retail. The same four routes at the Ofgem default-tariff cap in force on each date, which is a step function of policy decisions rather than a market series. The geothermal network route is priced at the **non-domestic** electricity rate, because a network operator buys on a commercial contract rather than the domestic cap. What is compared on that line is the operator's input cost, not a connected household's bill, which is a network tariff this model does not carry.

Three services, from v8.34. The chart carries a second toggle, and **both bases carry all three** — six combinations, not two plus three:

Service	Flow temperature	Gas efficiency
Space heating	Weather-compensated, 50 °C at -5 falling to 30 °C at +15	0.835
Hot water	52 °C year-round (†) — a cylinder needs its temperature whatever the weather	0.73
As delivered	Both, blended by the mix served that day	Both, blended

The blend is a harmonic mean weighted by delivered service, because what is being averaged is a cost per unit delivered:

$$\text{effective COP} = 1 / (w_{\text{space}}/\text{COP}_{\text{space}} + w_{\text{dhw}}/\text{COP}_{\text{dhw}})$$

with w_{space} taken from the same HDD-shaped, summer-damped daily weighting the hourly engine uses (Section 10.2), and hot water flat. An arithmetic blend would flatter whichever service is dearer. **This is the only one of the three that moves with the load**, and it is where the separate hot-water efficiency becomes visible: it tracks space heating through the winter and hot water through the summer, when almost nothing else is being asked for.

The hot-water view sets an honest limit on the argument. Because the cylinder pins the flow temperature, none of the flow-reduction benefit of Section 10.9 is available to it and the lift never shrinks. The geothermal network's advantage over a gas boiler is about **1.7× in both January and July** on hot water, against 1.9× rising to 4.4× on space heating. The case for heat pumps is strongest exactly where the flow temperature can fall, and weakest on the load that pins it — which is the same mechanism as the single-header problem in non-domestic buildings.

A service-volume band is drawn beneath the price chart, on the same window and the same x axis: delivered GWh per day, space heating stacked with hot water, the selected service

solid and the other faded. A price is only as interesting as the service it prices, and without the volumes a reader cannot weigh the three views against each other. By July the load is entirely hot water, so a space-heating price that month describes something nobody is buying. Space heat is shaped by the same HDD-weighted, summer-damped daily weighting used for the blend; hot water is flat at the annual rate.

On the retail basis the three services are built the same way, from the cap-derived unit prices divided by each service's own COP. Gas requires one extra step: it arrives already divided by the space-heating efficiency, so the raw unit price is recovered before being redivided at 0.73 or at the blend.

Correction, 14 August 2026. Between v8.34.0 and v8.34.1 the retail basis carried space heating only, and the service selector short-circuited on basis before it considered service — so on that basis the three buttons appeared to work and did nothing. No incorrect figure was displayed; the selection was simply ignored. Both are fixed, and the six combinations are now asserted in the front-end test.

Two electricity means are published per day: the flat mean, which the headline spark-gap ticker uses, and the **heat-weighted mean**, which is what a heat pump would actually pay because its load concentrates on the morning and evening peaks. The difference is the coincidence premium (Section 10.7), approximately 3% across a winter. The flat basis therefore understates every heat-pump route by about that much, and the site says so.

System Average Price is a balancing price — what National Gas pays to balance the system on the day. It tracks NBP day-ahead closely but is not the same instrument.

A caution on averaging. Any average of this series taken over *days* rather than over *heat delivered* will flatter whichever route performs best in mild weather. Unweighted, air-source shows a higher mean COP than ground source (4.88 against 4.06) and inverts the site's own SPF ordering; weighted by heat, the ordering is correct (3.49 against 3.66). On the thirty highest-heat days the figures are 2.21 and 3.08. Both weighted and unweighted values are published so the distinction is checkable.

9.4 What the value-for-money appraisal is, and is not

Added at v8.42 after peer review. Section 6 of the site was previously headed “How geothermal would improve the UK's heat networks” and its frontispiece claim read “making UK heat networks geothermal pays for itself, on the government's own appraisal rules”. Both overstated what the calculation establishes, and both have been changed.

It is one counterfactual. The comparison is geothermal against an air-source-led network with separate air-cooled chillers. That is a real and increasingly likely alternative, because zoning designates networks without settling their source. It is not the only one: a network can be fed by industrial waste heat, river or sewer heat, data-centre rejection, large water-source heat pumps, energy from waste, or a hybrid. **Nothing in the appraisal establishes that geothermal is the preferred national source**, only how it compares with one named alternative.

It is not a full Green Book appraisal. It applies Green Book conventions — the sixty-year horizon, the declining social time preference rate, DESNZ carbon values, optimism bias, switching values — to the *increment* between two sources. It does not carry a full options appraisal across the feasible source set, a distributional analysis, or the unmonetisable effects the Green Book also requires. The heading is now “An illustrative appraisal of geothermal against an air-source-led network” and the panel label reads “One counterfactual, on Green Book conventions — not a full appraisal”.

One load-bearing assumption remains stated on the panel: that distribution and connections are common to both routes and therefore cancel, which is approximately true for a given network design and not axiomatically true.

9.5 Pre-FID development capital

Brought inside the appraisal at v8.43, after peer review. It was previously excluded on the ground that a developer carries it at a development return rather than infrastructure capital at financial close. **That is a financing argument.** A social cost-benefit analysis counts a resource cost whoever bears it, and must count the prospects that never reached close.

The figure is **7.40% of the subsurface increment**, blended across the two things this scenario actually is:

Component	Devex	Success	Social	Weight
Shallow UTES / closed loop	4%	~100%	4.00%	60%
Intermediate doublet	10%	80%	12.50%	40%
Blended			7.40%	

Dividing by the success rate is the substantive step. Money spent on prospects that did not proceed is still a resource cost per delivered megawatt, and it is precisely what a developer’s committed-capex figure omits. A closed loop is *installed* rather than found and carries no resource risk — its development cost is design, permitting and thermal response testing. A doublet carries real risk.

The 10% is scenario-wide, not project-specific (†). Causeway’s own hospital case runs at 15%; two further hospital projects are de-risked by legacy oil and gas data, which is a UK-specific asset rather than a general condition.

Against the published benchmarks. Rystad’s geothermal economics model puts pre-FID at about **2% of global geothermal capex**, against 47% subsurface and just over half surface facilities. Two adjustments are needed before that transfers: it counts only schemes reaching final investment decision, and it is dominated by large deep power projects, where European geothermal district heating averages about 10 MWth and fixed development cost is a larger share. Rystad well failure rates run from **under 10% in Germany and Hungary to 30% in the Netherlands**; UK maturity — 20 MWt of geothermal heat installed today — sits nearer the Dutch end, which is why 80% rather than 90% is used.

Effect. The benefit-cost ratio falls from **2.53 to 2.37**, about 6%, and £1.0bn is added to the cost before optimism bias. It is insensitive to the blend weight: at 20% intermediate the ratio is 2.39 and at 80% it is 2.27. Devex is under 1% of the programme's own capital, against a geothermal increment of about a ninth of it.

It is applied to the subsurface increment only, not to the cooling increment: cooling plant is specified, not explored for.

10.3 Coefficient of performance and calibration

Substantially revised at v8. The equations are given in full because the flow-temperature results in Section 10.9 depend on them.

The Carnot form

For a heat pump lifting heat from a source at T_s to a delivery temperature T_f , both in kelvin, the thermodynamic ceiling is the Carnot coefficient of performance:

$$\text{COP}_{\text{Carnot}} = T_f / (T_f - T_s) = T_f / \Delta T$$

Delivery temperature over lift. Real machines achieve a fraction of this:

$$\text{COP} = \eta \cdot T_f / \Delta T$$

with a floor of 8 K on ΔT (†), because the ratio diverges as the lift closes and no real machine does.

Route sources

Route	Source temperature
Air-source	Outdoor air less a 3 K evaporator approach, plus a defrost penalty of up to 12% between -5 and +7 °C
Ground source	8.0 °C — 11 °C ground less a 3 K brine approach
Geothermal network	19.6 °C (†) — a charged ambient loop, UTES or intermediate-doublet

The network source is 19.6 °C, not the 10 °C stated in v7.1. The correction is material: it is the temperature a *charged* loop reaches, and the difference between it and virgin ground is the basis of the seasonal-storage argument in Section 10.9.

Flow temperature

Flow is weather-compensated linearly between endpoints:

$$T_f(T_{\text{out}}) = 50 \text{ °C at } T_{\text{out}} = -5 \text{ °C, falling to } 30 \text{ °C at } T_{\text{out}} = +15 \text{ °C,}$$

clamped outside that range (†)

Calibration of η

η is **not assumed**. For each route it is solved so that the annual heat-weighted mean of the hourly COP reproduces that route's published seasonal performance factor. The current calibrated values are:

Route	η	Anchor SPF
Air-source	0.33198	2.80
Ground source	0.33646	3.24
Geothermal network	0.33436	5.00

The agreement of the three fractions is NOT independent evidence, and v8.36 withdrew the claim that it was. For air-source and ground source both the source temperature and the SPF are sourced, so η is genuinely solved and its value is informative. For the network route neither is: the SPF of 5.0 is an estimate, and the source temperature of 19.6 °C was **chosen so that the network reproduces that SPF at the same Carnot fraction as the other two**. The match is therefore how the source was derived, not a test of it.

Earlier versions of this document described the agreement as “a check, not a coincidence”. That was wrong, and the code comment in `retro.py` said so at the time — the constant is annotated “DERIVED, not measured”. The claim has been removed from the methodology and from the site.

What the pair does and does not establish. It establishes that SPF 5.0 and a 19.6 °C source are mutually consistent under the same Carnot fraction that fits two monitored fleets, which is a meaningful plausibility statement and not a trivial one. It does not establish either quantity independently, and a reader should treat them as one estimate with two faces rather than two estimates that agree.

Sensitivity. Because the two are linked, the honest presentation is a matrix rather than a point. At the calibrated fraction, a network source of 15 °C supports an SPF of about 4.3 and 25 °C supports about 6.2; conversely, holding the source at 19.6 °C, an SPF of 4.0 implies a Carnot fraction of 0.27 and an SPF of 6.0 implies 0.40 — the first below and the second above anything in the monitored fleets. That range, rather than the point value, is what the evidence supports.

The 15% drift gate is retained as a **wiring check** — it catches a route whose calibration has failed to converge — and is not a validation.




Figure 8

Figure 8. (a) The weather-compensation curve and (b) the point COP it produces for each route. Air-source falls toward 2 in the cold, exactly when heat is most wanted; a ground loop barely moves.

10.9 Sensitivity to flow temperature

New at v8.

Differentiating the Carnot form with respect to flow temperature, and noting that both numerator and denominator depend on T_f :

$$\partial \text{COP} / \partial T_f = \eta \cdot (-T_s) / \Delta T^2$$

The flow temperature cancels from the numerator; only the source survives. As a fractional change — which is what “per cent per kelvin” means:

$$(1/\text{COP}) \cdot \partial \text{COP} / \partial T_f = -T_s / (T_f \cdot \Delta T)$$

η cancels, so the percentage gain is independent of machine quality, and **ΔT remains in the denominator**, so the gain per kelvin is inversely proportional to the lift. Halving the lift doubles the value of each further kelvin (Figure 9).




Figure 9

Figure 9. Fractional COP gain per kelvin of flow reduction, against lift. The route with the smallest lift gains most. The shaded band is the HeatpumpMonitor field range of 2.7 to 3.1 per cent per kelvin.

At a 45 °C flow: ground source at an 8 °C source has $\Delta T = 37$ K and gains 2.4% per kelvin; the network at 19.6 °C has $\Delta T = 25.4$ K and gains 3.6%. The ratio is simply $37/25.4$.

Comparison with field data. HeatpumpMonitor.org publishes, from several hundred MID-metered UK systems at the SEPemo H4 boundary, the regression

$$\text{SPF}_{\text{H4}} = -0.1054 \cdot (\text{heat-weighted flow} - \text{outside}) + 6.8762, \quad R^2 \ 0.547$$

This is the strongest correlation in that dataset. It implies **2.7–3.1% per kelvin** over a 25–35 K lift, against this model’s 1.9% for air-source at the same condition (Figure 9). **The model therefore understates the effect on the one route for which field data exists.**

The most likely cause is that η is held constant. If the Carnot fraction itself improves as the lift shrinks — less superheat, better part-load matching — the field slope contains two effects where this model contains one. That would mean the network figure is conservative rather than optimistic. It is being put to OpenEnergyMonitor for comment and should be regarded as open.

The network figure of 3.6–4.2% per kelvin extrapolates beyond any monitored system and is daggered accordingly.

10.5 The capacity line

Changed at v8.

Requirements are set against a capacity line constructed as a de-rated block of dispatchable and interconnector capacity **plus the wind and solar actually generated in each hour**. Because outturn wind and solar are added hourly, the block must **exclude** de-rated wind, or wind is counted twice.

The block is derived identically for each winter from the NESO Winter Outlook:

block = ACS peak demand + de-rated margin – de-rated wind

Winter	ACS peak	Margin	Total de-rated	Less wind	Block
2024/25	59.8 GW	5.2	65.0	4.4	60.6 GW
2025/26	61.0 GW	6.1	67.1	~4.8 †	62.3 GW

The 2024/25 wind de-rating is NESO’s own (30,532 MW installed de-rated to 4,416 MW). The 2025/26 figure is not published in the summaries and is scaled with installed capacity growth (†); the margin is sourced, the wind split is estimated.

Each winter is tested against its own fleet. Prior to v8 a single 62.0 GW block was applied to every hour, which was defensible while the store held one winter and became wrong the moment it held two.

One consequence for the binding-hour definition. With a ceiling that varies by winter, the binding hour is the hour of **least headroom**, not of highest absolute requirement. The two coincide when the store holds a single winter and need not otherwise. Both `_elec_limits` and `system_view` therefore maximise (requirement – block) and publish the block each figure was measured against.

11. Automation, failure handling and reproducibility

(revised at v8: gates added)

A scheduled GitHub Actions workflow executes daily at 03:43 UTC. In addition to the four hourly gates (Section 10.8), three whole-build gates were added at v8 following defects that reached publication:

Annual-quantity plausibility. Trailing-twelve-month heating degree days must fall within 1,200–3,200. This gate exists because raising the regression window from 365 to 730 days caused every quantity that treated the window as a year to double — `hdd_trailing_12m` went 2,017 to 4,114, calibrated annual gas space heat 261.8 to 531.5 TWh, and the electrification limit halved from 98% to 49% — while all four test suites passed, because they check structure and internal consistency rather than whether an annual quantity spans a year. The window and the year are now separate constants (`WINDOW_DAYS = 730`, `ANNUAL_DAYS = 365`) and every trailing-twelve-month sum slices the latter.

Document structure. The published HTML is checked for content before the doctype, balanced script tags, and JavaScript leaking into markup. A failed string replacement once

wrote 1,341 characters of JavaScript in front of the doctype, which rendered as visible text at the top of the live page while every syntax check passed: node - - check on an extracted script proves the script parses, not that the document is well formed.

Scope. The build script's `main()` is walked for names read but never assigned in that scope and not module-level. A syntax check cannot see an out-of-scope name.

12. Principal limitations

(revised at v8)

The following are the material limitations, restated for peer review.

— **The ECUK cooling and ventilation split** is now sourced (54:46, Section 6.1) but rests on a 2014–15 survey applied to later metered totals. No comprehensive UK non-domestic end-use survey has superseded BEES.

— **The cooling EER of 3.0** (range 2.5–4.0) propagates one-for-one into every cooling service figure and is the largest single unverified assumption in that chain (Section 5.4).

— **The comfort-deficit stock fractions and thermal-response coefficients** carry a several-fold range (Section 5.2).

— **The observed cooling response is a lower bound**, not an independent estimate, because within-class centring removes the level (Section 4.5). The undemeaned reconciliation now recovers 75% of the anchor on its base specification and 95% on the full one (Section 4.6), but is not yet used in the published figures.

— **The geothermal network SCOP of 5.0 and source of 19.6 °C** have no monitored UK fleet behind them. Every network figure on the site inherits this.

— **The flow-temperature sensitivity for the network route** extrapolates beyond any monitored system, and the model is known to understate the effect on the one route where field data exists (Section 10.9).

— **Geothermal output is not nationally metered**; the present-day bar remains the most weakly constrained figure on the site.

— **The capacity line is a static overlay**, now on **two** winters of observed weather rather than one, with per-winter blocks that must be revised each autumn. It cannot represent storage duration or network constraints.

— **Two winters is still a small sample.** The binding hour is now chosen across winters rather than describing the only one available, and 2025/26 holds it against 2024/25's own lower ceiling — but two is not a distribution.

— **Annual anchors lag** by roughly nine to fifteen months.

— **Hourly heat shape is modelled from weather, not metered.**

13. Corrections declared

New at v8. The site's practice is to declare corrections rather than silently restate. The following were found and fixed between v7.1 and v8, and are listed because a reviewer should be able to see what the process catches.

Defect	Effect	Resolution
Regression window used as if it were a year	Every annual quantity doubled; electrification limit halved 98% → 49% for one day	ANNUAL_DAYS separated from WINDOW_DAYS; plausibility gate added
Comfort-deficit threshold applied an indoor criterion to outdoor air	Tier understated roughly five-fold	Rebased to 20 °C daily mean from EFUS
Binned response curve imposed a ceiling	Hot days understated; annual total 26% low	Continuous quadratic fit
Saturation tier premise contradicted by the data	A published band that was structurally zero	Tier withdrawn; computation kept as a diagnostic
Ventilation credited with a coefficient of performance	Cooling service overstated by about one third	Ventilation reported separately at unity
Single capacity block applied to two winters	2024/25 tested against 2025/26's fleet	Per-winter blocks
Network source stated as 10 °C in v7.1 methodology	Understated the network COP and omitted the storage mechanism	Corrected to 19.6 °C, charged loop
One boiler efficiency applied to both end uses	Hot water credited with a condensing recovery it does not achieve; gas had no seasonal efficiency shape	Separate SAP-derived hot-water efficiencies, gas and oil
Humidity null obtained under a bias running toward null	The original null rested on the contaminated variable; two corrected partial-record reruns then reversed it twice	Rerun on the full 800 days with solar carried as a control: dewpoint +0.0006, enthalpy +0.0000, quadratic moves 3%. Null properly obtained; extensive margin re-established
Cooling-fit workflow input overrode the fetch default	The humidity fetch ran at 400 days against an 800-day store, making moisture terms incomparable with every other control	Workflow default raised to 800, stale series cleared before fetching, and the run now fails below 700 days

Defect	Effect	Resolution
Embedded solar added twice in three diagnostics	The diurnal comfort/process separation was an artefact (2.59x collapses to 0.93); the 74% cooling reconciliation invalid. Night-hour results unaffected, solar being zero	All three call sites corrected; diurnal chart and 74% figure withdrawn; found in peer review
Network source presented as an independent check on the Carnot fraction	19.6 °C is derived to make the fractions agree, so the agreement evidenced nothing	Claim withdrawn from methodology and site; sensitivity range published instead
Front-end added the air-source approach instead of subtracting it	Displayed air-source COP too favourable - conservative against the site's own case, but wrong	Sign corrected; fallback given the same sign as the payload
"Electrification limit" implied capacity-adequacy modelling	The calculation is a historical headroom stress test and carries no network, outage or storage modelling	Renamed; scope stated on-panel
Pre-FID development capital excluded from a social appraisal	A financing argument used to exclude a resource cost, and prospects that never reached close were uncounted	Brought inside at a blended 7.40% of the subsurface increment; BCR 2.53 to 2.37
VFM headline claimed a general heat-network result from one counterfactual	"Making UK heat networks geothermal pays for itself on the government's own appraisal rules" overstated an illustrative single-counterfactual appraisal	Heading, label and point 06 claim rescoped; scope statement added to the panel
Network SPF 5 system boundary not stated	A whole-system H4 figure and a plant-only COP are not comparable, and the reader could not tell which was meant	Boundary listed explicitly, inside and outside; parasitic sensitivity published
Weekly Tier 1 kept the binned estimator after the daily series moved to a continuous fit	Same metric, two implementations; the weekly bar understated a hot week by about 80%	Fit lifted above both accumulations; one estimator; found in peer review
Gas-network GB factor applied to every fuel	GB oil heat overstated by about 19%; total GB annual heat about 4% high	Per-fuel shares from NISRA household data; found in peer review
Tier 1's share of the anchor stated as "about a seventh"	The published series had moved to a continuous quadratic and other	Recomputed from the payload and corrected in three places

Defect	Effect	Resolution
	corrections; the true share is about a third, so the site understated its own measured response by a factor of 2.3	
Price-chart service selector short-circuited on basis	On the retail basis the service toggle was ignored; space heating shown whatever was pressed	Six explicit basis-service combinations; retail series extended to all three

14. Change log against v7.1

Section	Change
4.3	Binned response replaced by a continuous quadratic; latent-slope extrapolation withdrawn
4.5	Observed-versus-anchored discrepancy explained rather than noted
4.6	New — undemeaned baseline model, balance point identification
4.7	New — diurnal separation and night elbow
4.8	New — humidity tested; result withdrawn at v8.37 pending a rerun on corrected data
5.1	Threshold rebased 26 °C indoor – 20 °C outdoor daily mean
5.3	Three tiers reduced to two; Tier 0 named
5.4	Cooling EER 3.5 – 3.0 with a published range
4.9	New — the cooling floor: four routes tried, two withdrawn, one open
5.5	New — per-fuel UK-to-GB shares
6.1	New — BEES 54:46 split; ventilation excluded from cooling; base 18.0 – 15.5 °C; separate hot-water boiler efficiencies from SAP Table 4b
6.4	New — heat-price series, two bases x three services, with a delivered-volume band
10.3	Carnot equations given in full; η calibration table published; network source 10 – 19.6 °C
10.5	Per-winter capacity blocks; binding hour redefined as least headroom
10.9	New — flow-temperature sensitivity and the

Section	Change
	field comparison
11	Three whole-build gates added
12	Limitations restated; one-winter caveat updated to two
13	New — corrections declared